

Foundations of Differentially Oblivious Algorithms

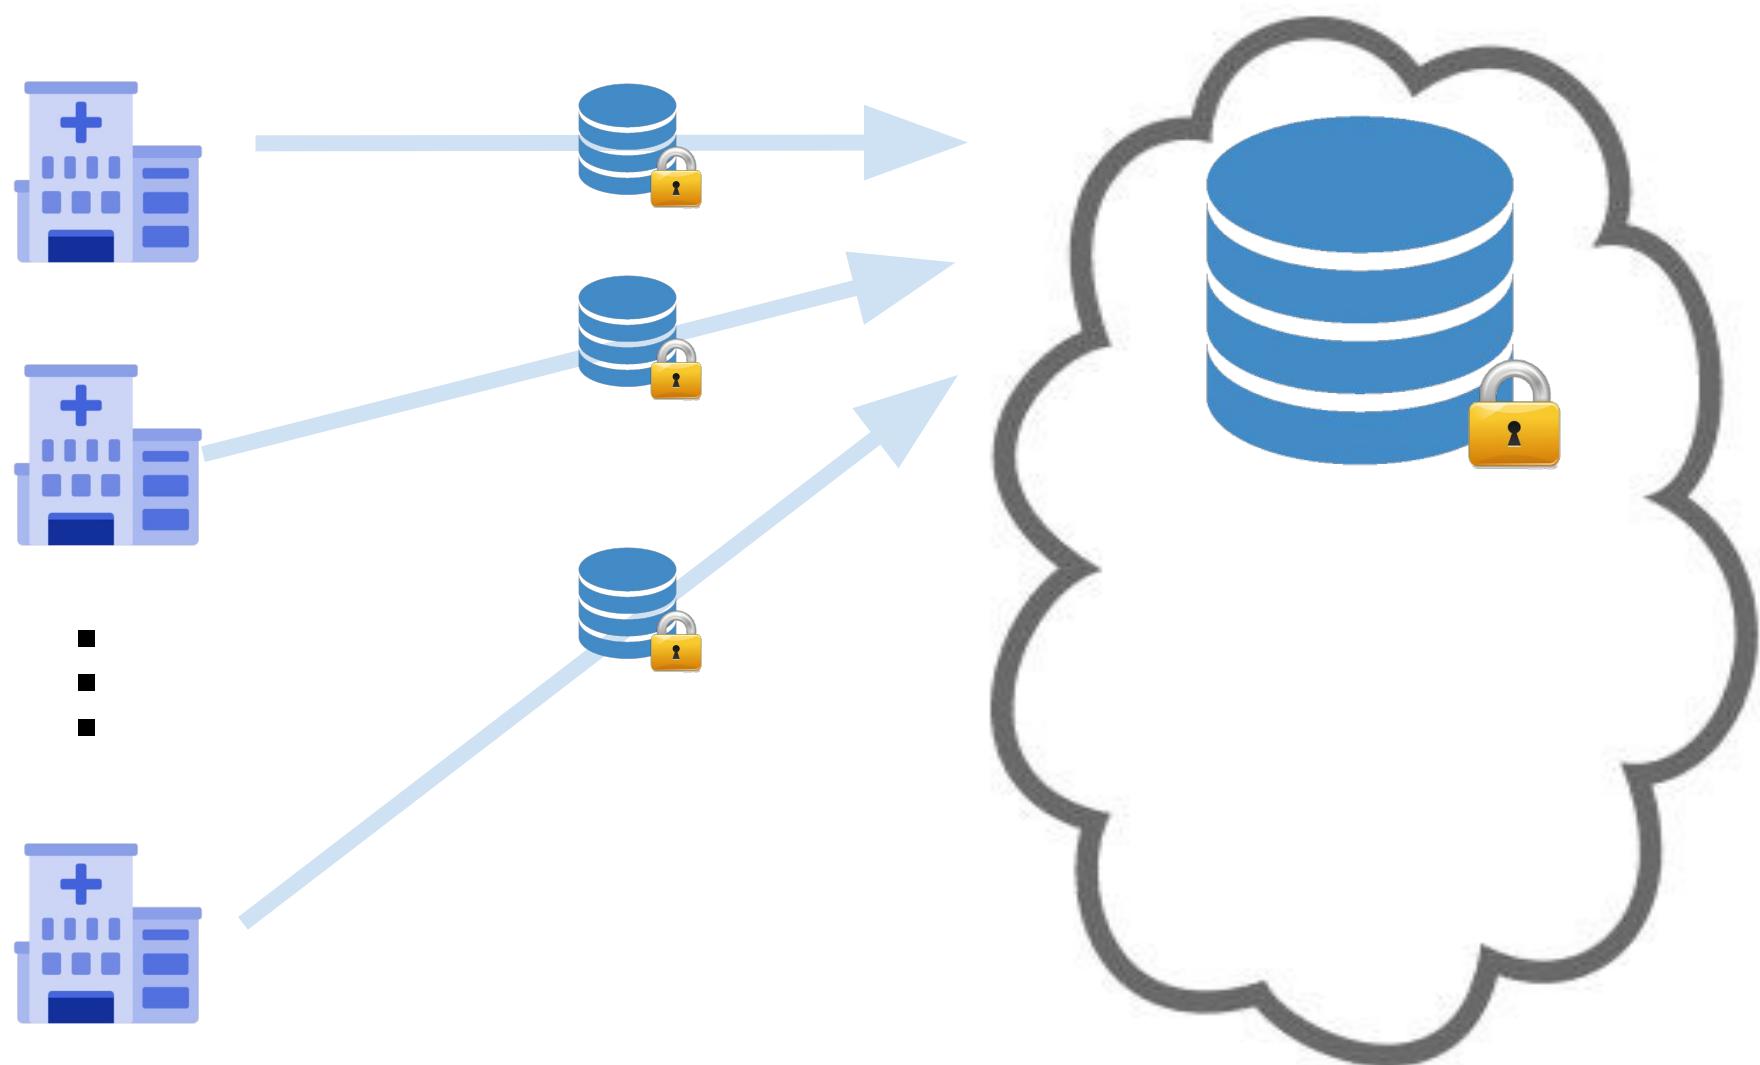
Elaine Shi

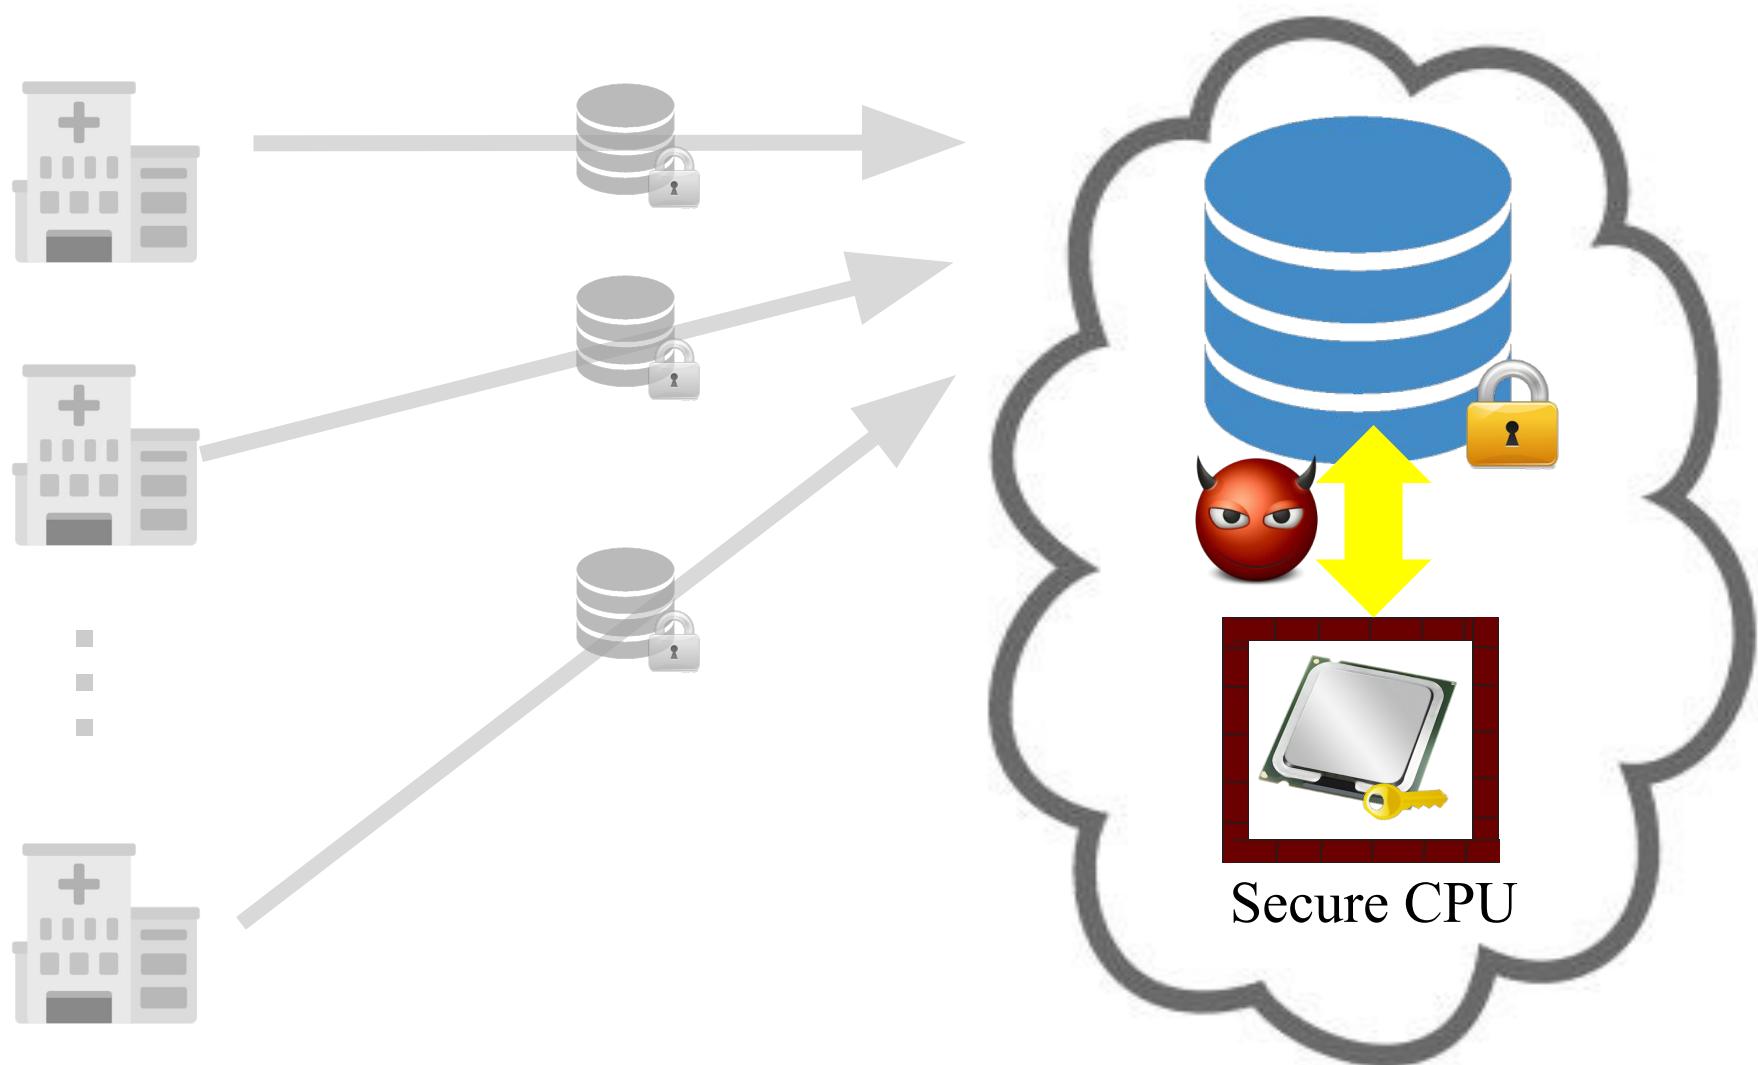
Based on [CCMS'18] and [LSX'18]



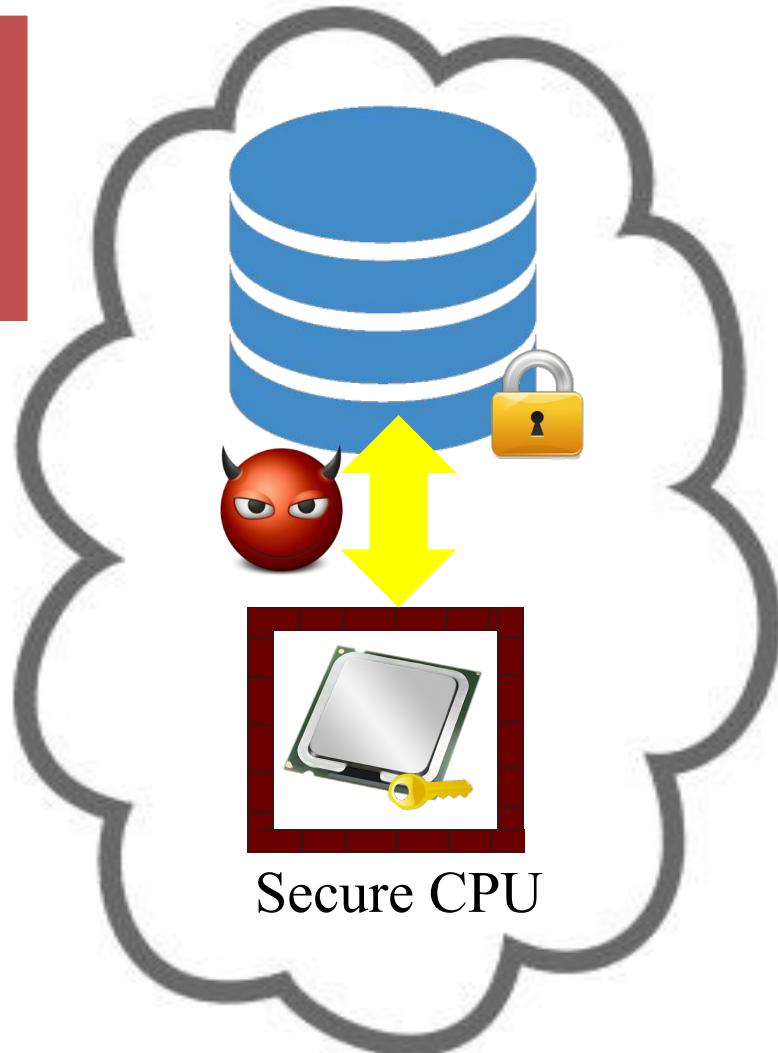
⋮







Access patterns to even
encrypted data leak sensitive
information.



Access Pattern Attack: Computing on JPEG Image

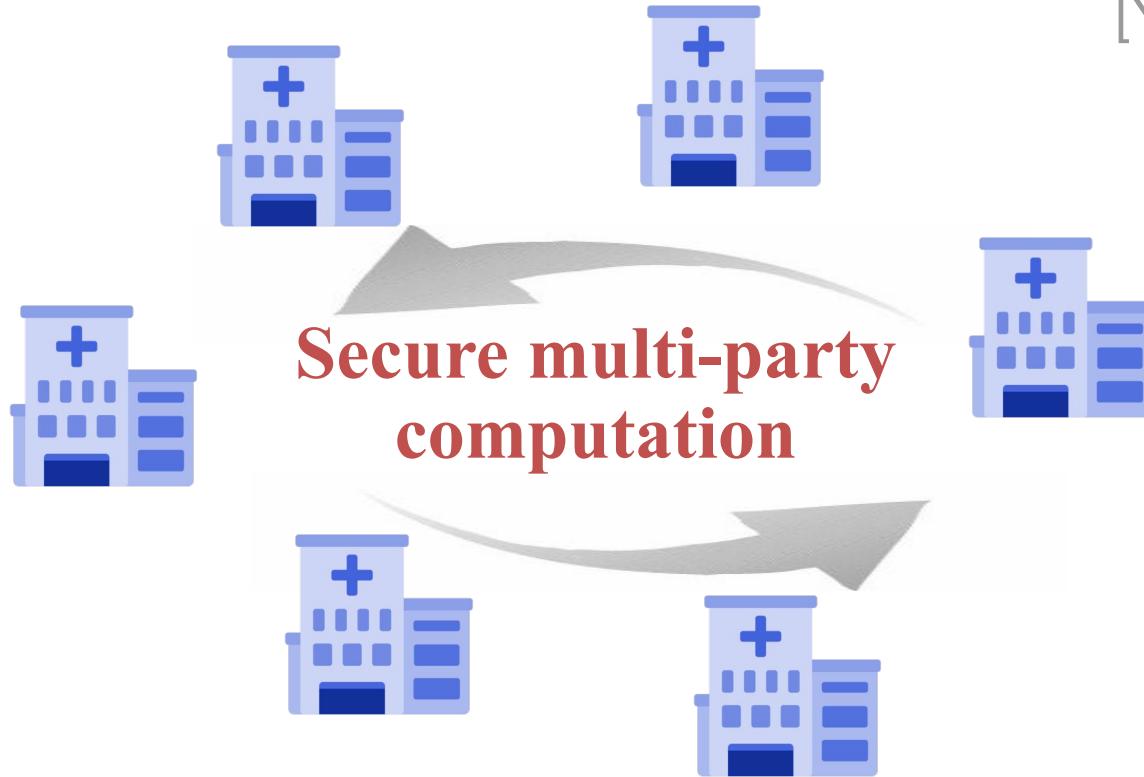
Original



Recovered



[Yao'82, GMW'87]



Access Pattern Leakage in MPC

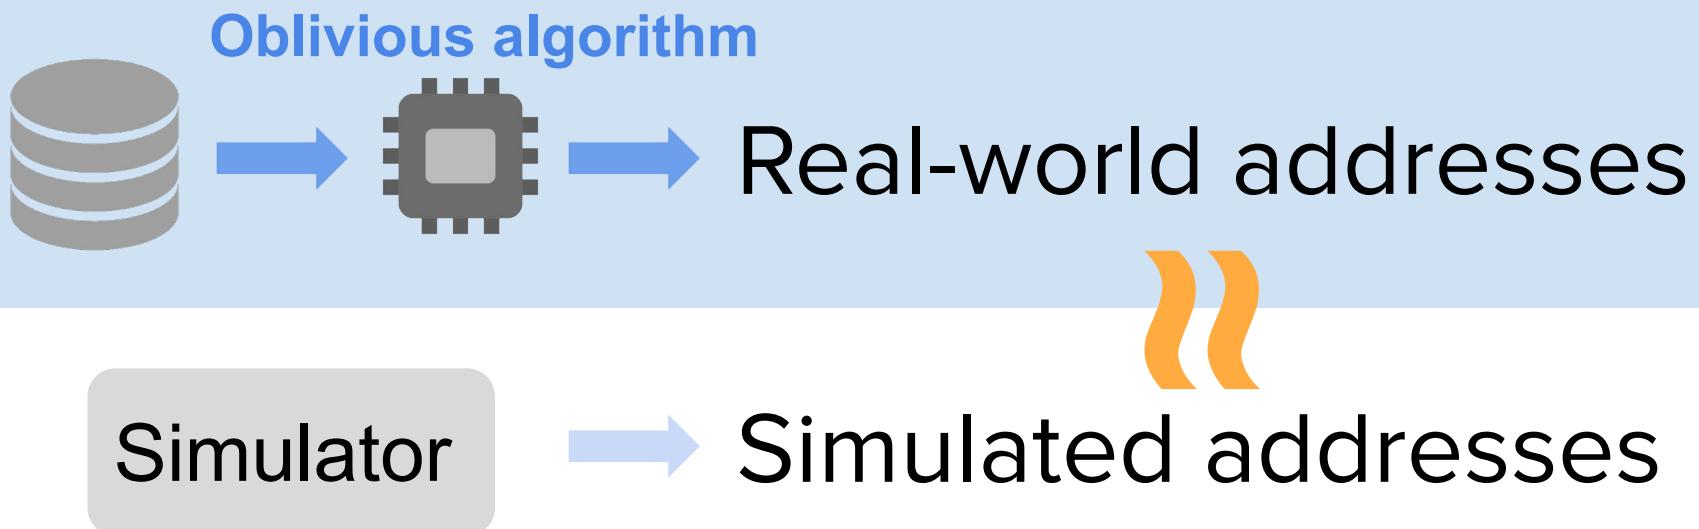


Oblivious RAM

An algorithmic approach that
provably obfuscates access patterns



Oblivious RAM





Oblivious RAM

“Encrypting the access patterns”



Oblivious RAM

“Encrypting the access patterns”

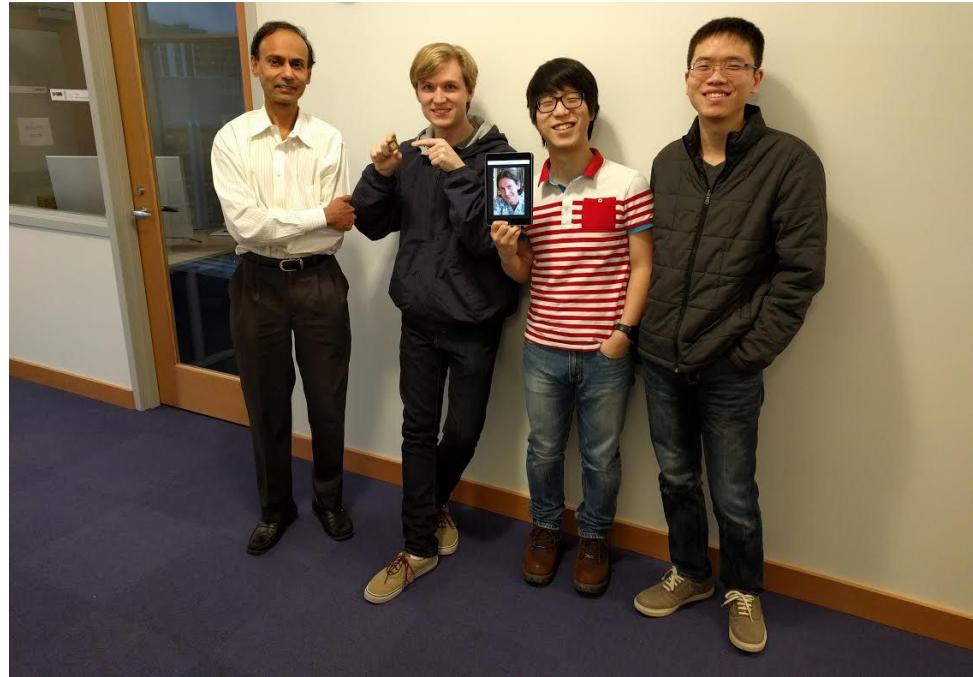
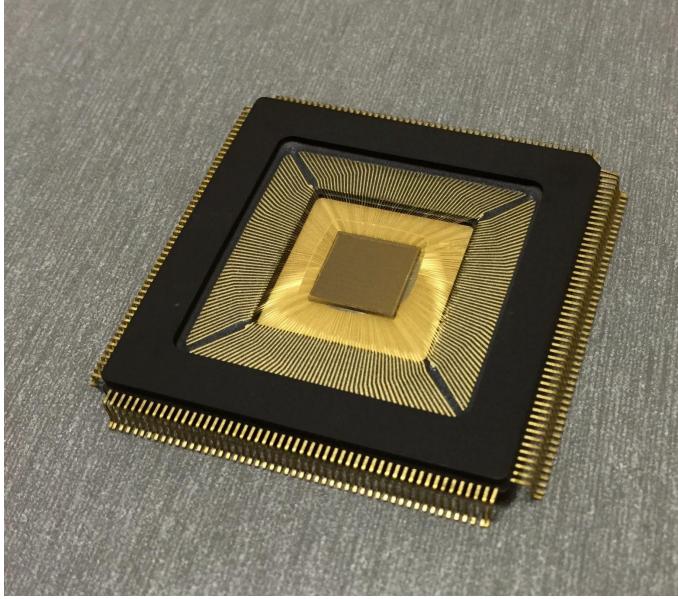
- Permute data in memory
- Shuffle data upon accesses

ORAM State of the Art



Any program can be made oblivious with
 $O(\log N)$ to $O(\log^2 N)$ overhead

[Optoroma, Circuit ORAM, ...]



ORAM State of the Art



Any program can be made oblivious with
 $O(\log N)$ to $O(\log^2 N)$ overhead



$\Omega(\log N)$ is necessary

[GO'96, LN'18]

ORAM State of the Art



Any program can be made oblivious with
 $O(\log N)$ to $O(\log^2 N)$ overhead

Implicit assumption:

Runtime is fixed and known



Implicit assumption:

Runtime is ~~fixed~~ and known

- Must pad to worst-case runtime
- Can incur even **linear** overhead



Implicit assumption:

Runtime is ~~fixed~~ and known

Relax the obliviousness notion?

- Still provide meaningful privacy
- Significantly improve efficiency

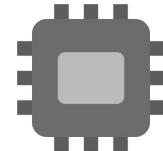
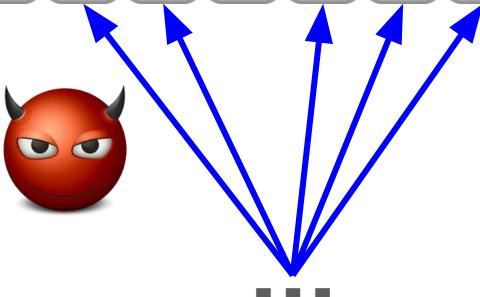




Differential Obliviousness

Inspired by differential privacy [Dwork et al. 05]

Memory



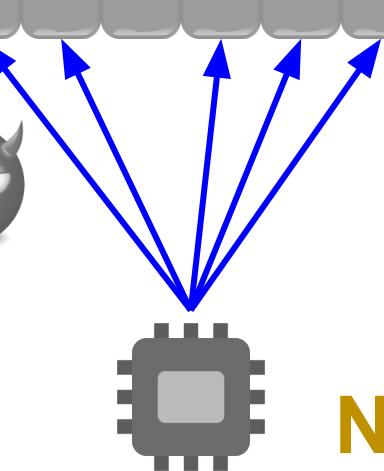
randomized

Algorithm
(e.g., compaction,
sorting)

Database



Memory



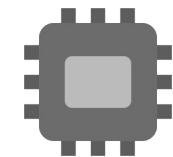
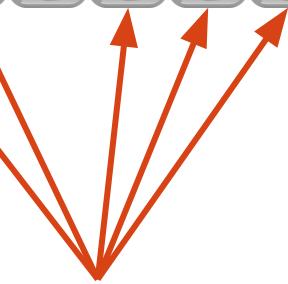
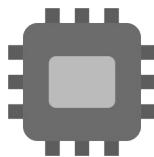
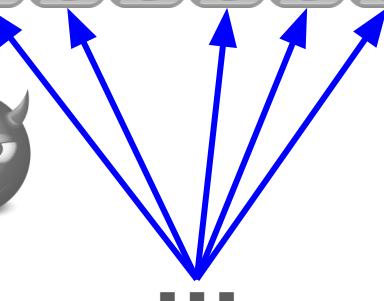
Algorithm
(e.g., data analytics)

Neighboring input DBs

Database



Memory

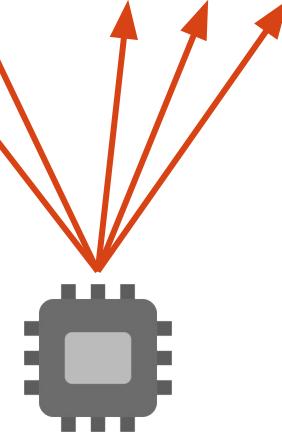
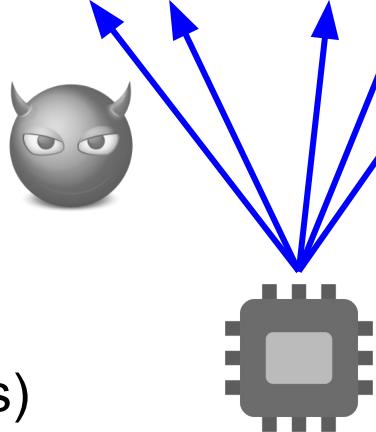


Algorithm
(e.g., data analytics)

Database

Access patterns on **neighboring** DBs must be **close**

Memory

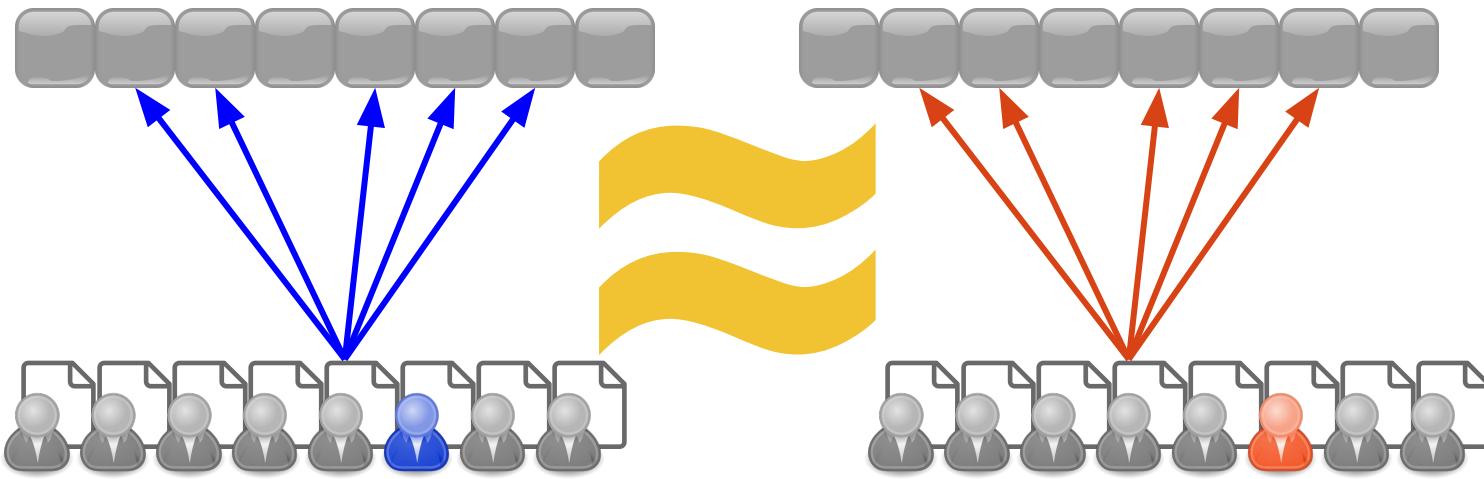


Algorithm
(e.g., data analytics)

Database



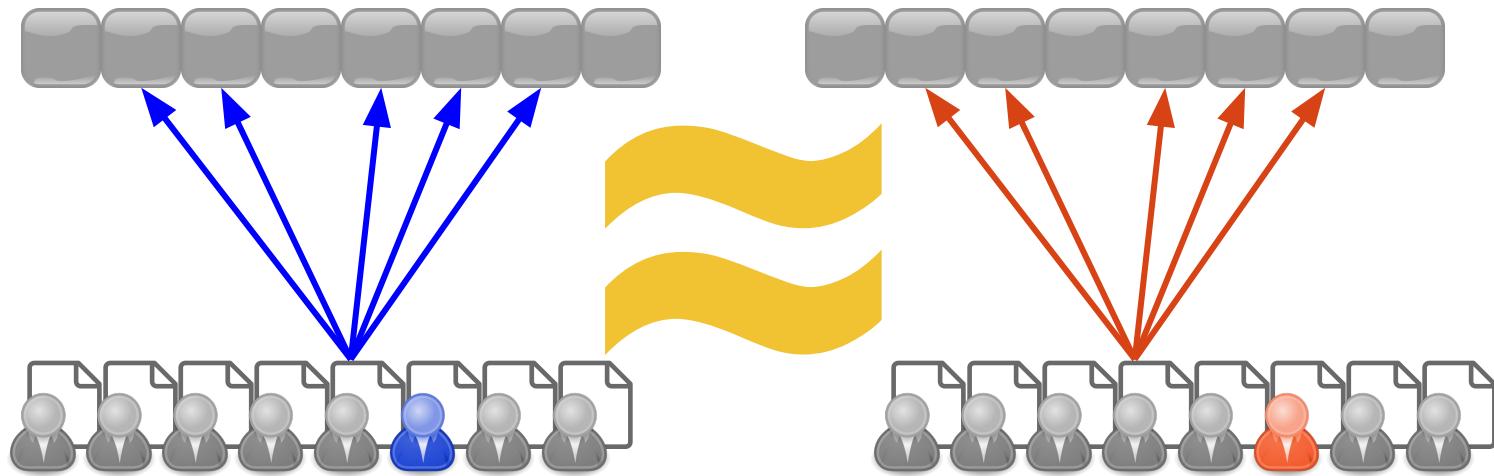
Access patterns on **neighboring** DBs must be **close**



$$\Pr[\text{blue arrows} \in S] \leq e^\epsilon \Pr[\text{orange arrows} \in S] + \delta$$

This must hold for any S

(ϵ, δ) -Differential Obliviousness



$$\Pr[\text{blue arrows} \in S] \leq e^\epsilon \Pr[\text{red arrows} \in S] + \delta$$

This must hold for any S

(ϵ, δ) -Differential Obliviousness

- 1 What is being relaxed?
- 2 Still provide meaningful privacy?
- 3 Overcome obliviousness barriers?

(ϵ, δ) -Differential Obliviousness

- 1 **What is being relaxed?**
- 2 **Still provide meaningful privacy?**
- 3 **Overcome obliviousness barriers?**

What is being relaxed?

Closeness needs to hold
only for neighboring DBs

$$\Pr[\text{blue arrows} \in S] \leq e^\epsilon \Pr[\text{red arrows} \in S] + \delta$$

This must hold for any S

What is being relaxed?

Closeness needs to hold only for neighboring DBs

Allow multiplicative,
non-negl. loss

$$\Pr[\text{blue arrows} \in S] \leq e^\epsilon \Pr[\text{red arrows} \in S] + \delta$$

This must hold for any S

Does not require padding
to worst-case runtime

$$\Pr[\text{blue arrows} \in S] \leq e^\epsilon \Pr[\text{red arrows} \in S] + \delta$$

This must hold for any S

(ϵ, δ) -Differential Obliviousness

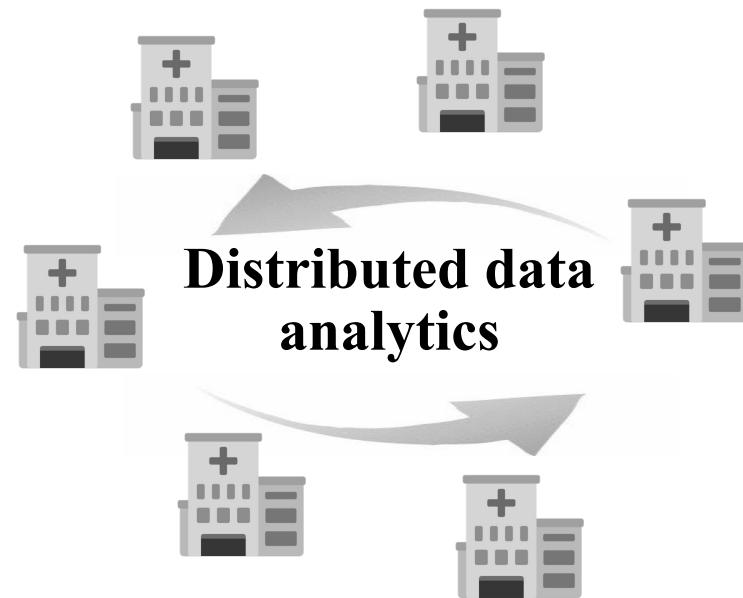
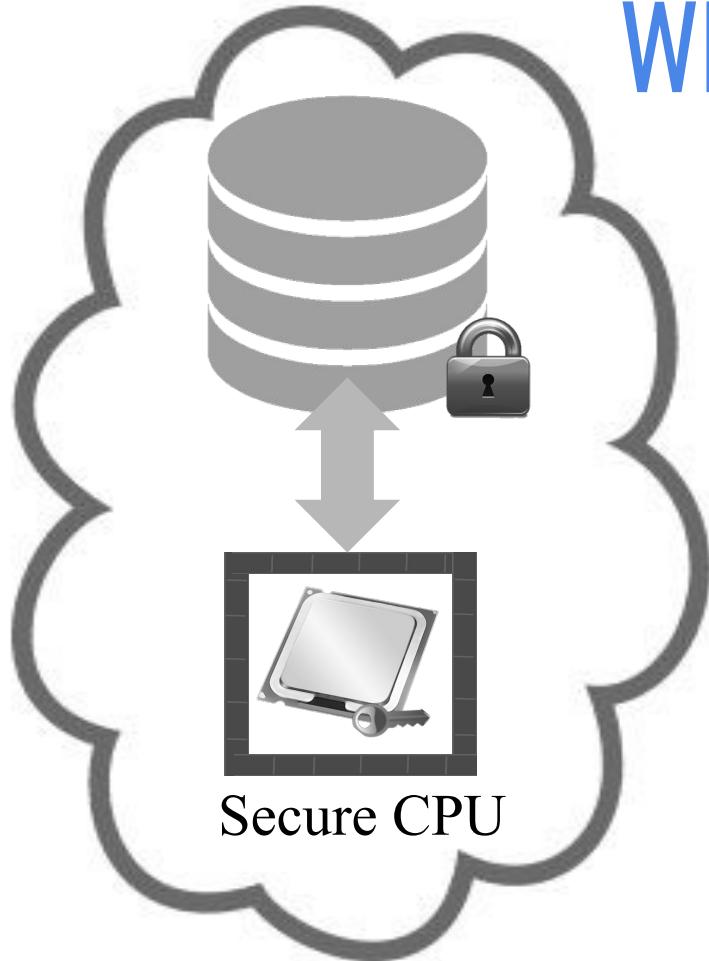
- 1 What is being relaxed?
- 2 Still provide meaningful privacy?
- 3 Overcome obliviousness barriers?



Bad idea if you are protecting your
Bitcoin signing key!

2 Still provide meaningful privacy?

When does DO make sense?



Typical parameters

Constant

Negl. in N

$$\Pr[\text{blue arrows} \in S] \leq e^\epsilon \Pr[\text{red arrows} \in S] + \delta$$

This must hold for any S

(ϵ, δ) -Differential Obliviousness

- 1 What is being relaxed?
- 2 Still provide meaningful privacy?
- 3 Overcome obliviousness barriers?

Stable Compaction

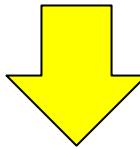
Obliviousness

$\Omega(N \log N)$
necessary

Differential
Obliviousness

$O(N \log \log N)$

Stable Compaction



Stable Compaction: Why do we care?

- Simple yet non-trivial
- Frequent algorithmic building block
- Warmup scheme in paper

Stable Compaction: **insecure** algorithm



Stable Compaction: **insecure** algorithm



Stable Compaction: **insecure** algorithm



Stable Compaction: **insecure** algorithm



Stable Compaction: **insecure** algorithm



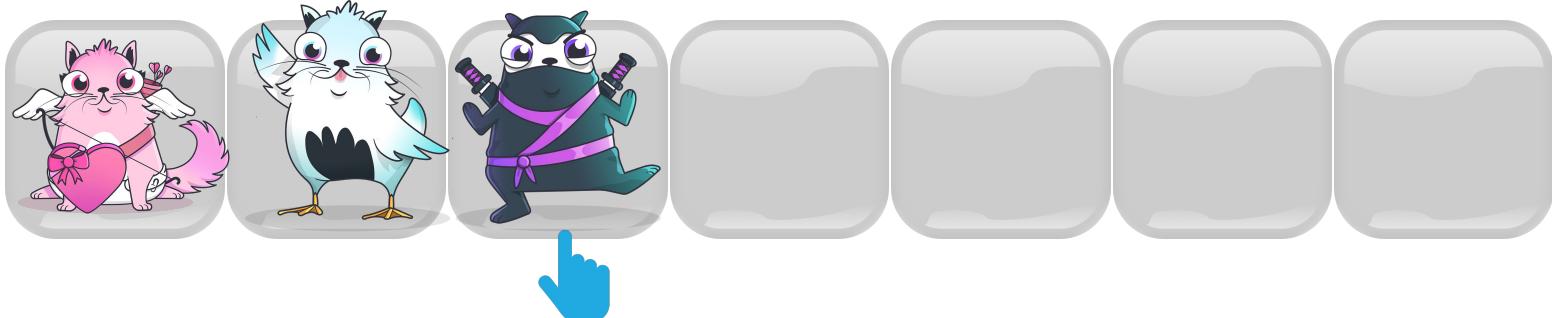
Stable Compaction: **insecure** algorithm



Stable Compaction: **insecure** algorithm



Stable Compaction: **insecure** algorithm



Stable Compaction: **insecure** algorithm

Completes in $O(N)$ time

Leaks exact progress

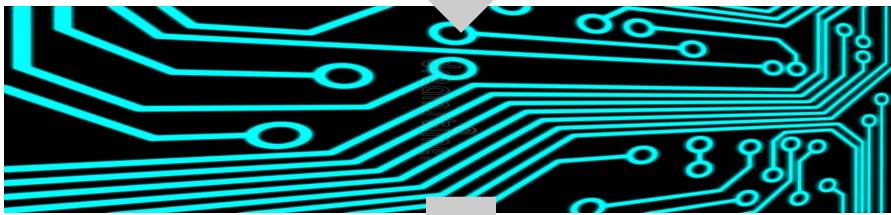
Stable Compaction: **oblivious** algorithm



Stable Compaction: **oblivious** algorithm



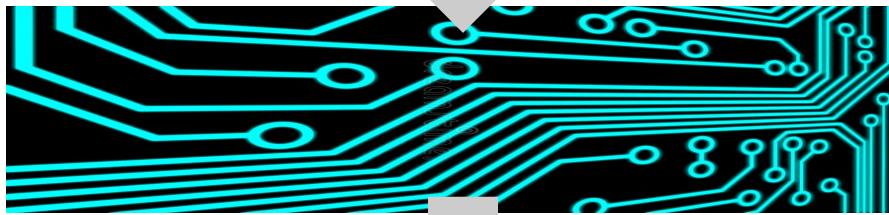
Sorting
network



Takes $N \log N$ time



Sorting
network

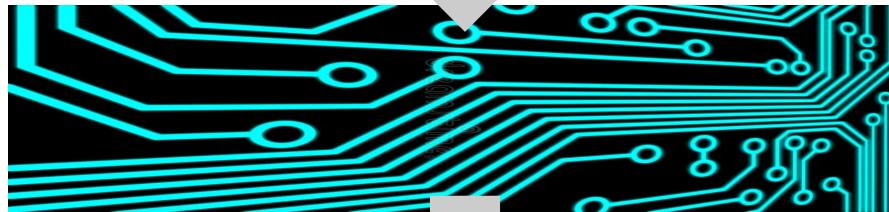


$N \log N$ time is necessary for obliviousness

Assumption: algorithm does not perform encoding on the kitties



Sorting
network



Stable Compaction

Obliviousness

$\Omega(N \log N)$
necessary

Differential
Obliviousness

$O(N \log \log N)$

Obliviousness

Cannot leak progress



Differential Obliviousness

Leak rough notion of
progress



polylog(N) batch

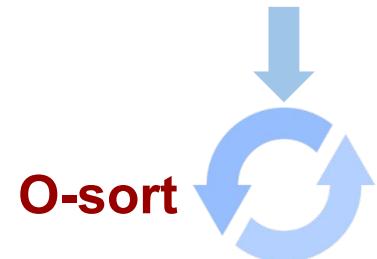
2~5 kitties so far



DP oracle



polylog(N) batch



2~5 kitties so far

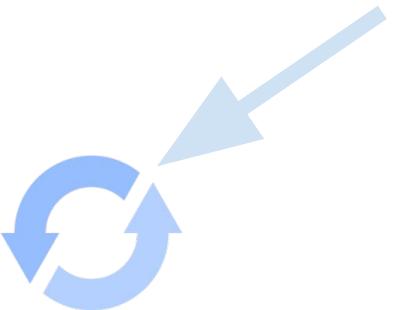
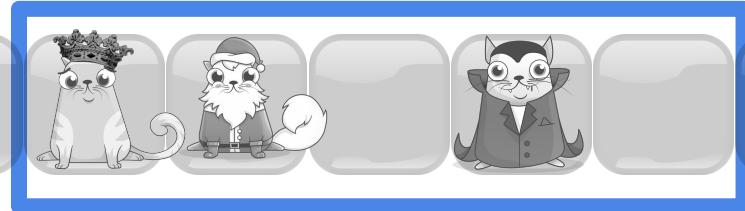


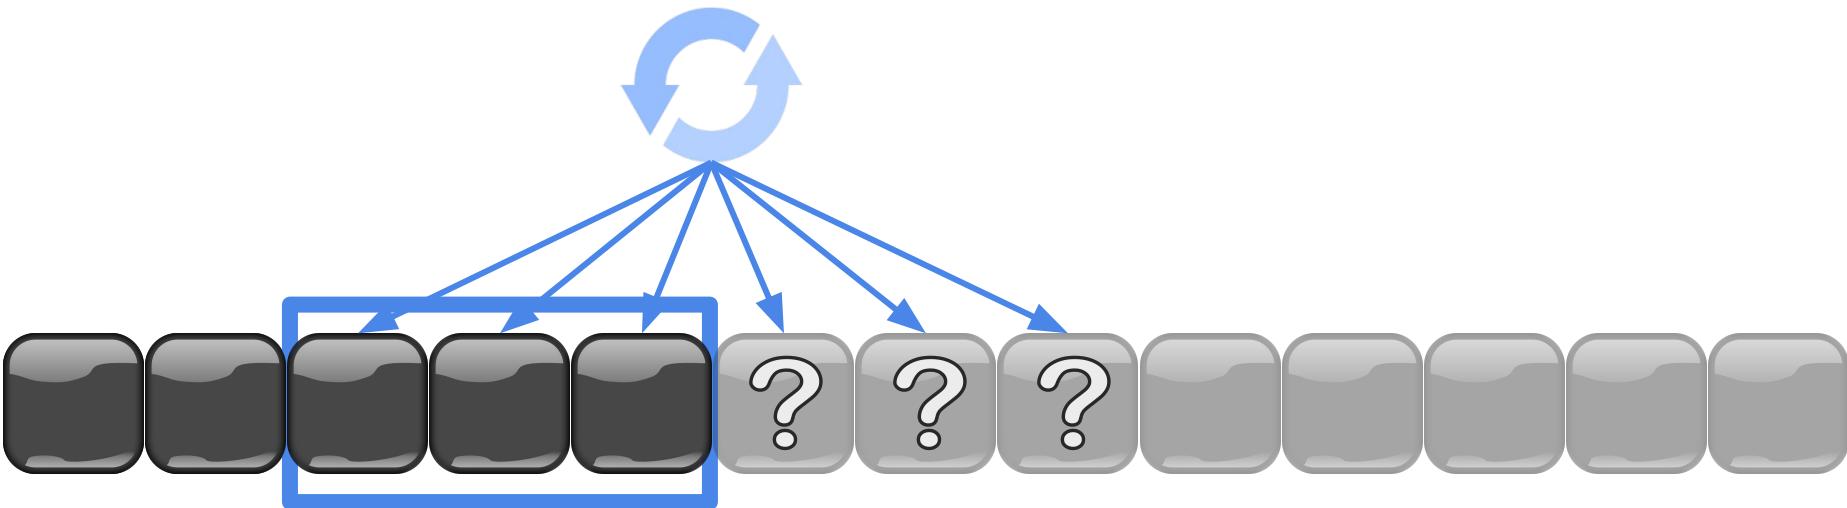
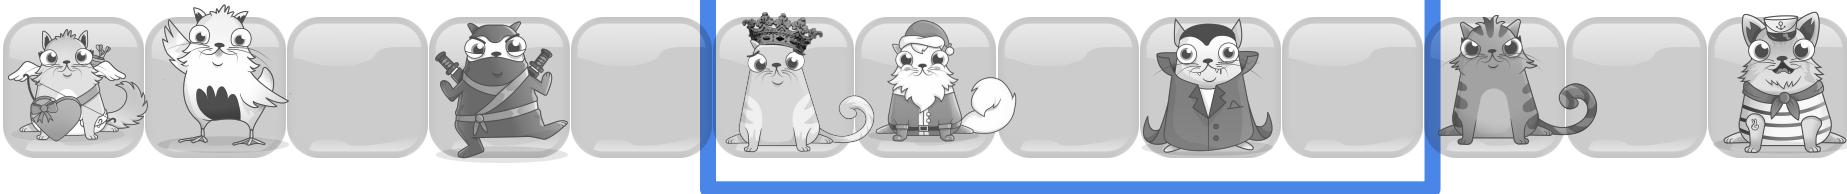


$\text{polylog}(N)$ batch



$\text{polylog}(N)$ error,
DP estimate

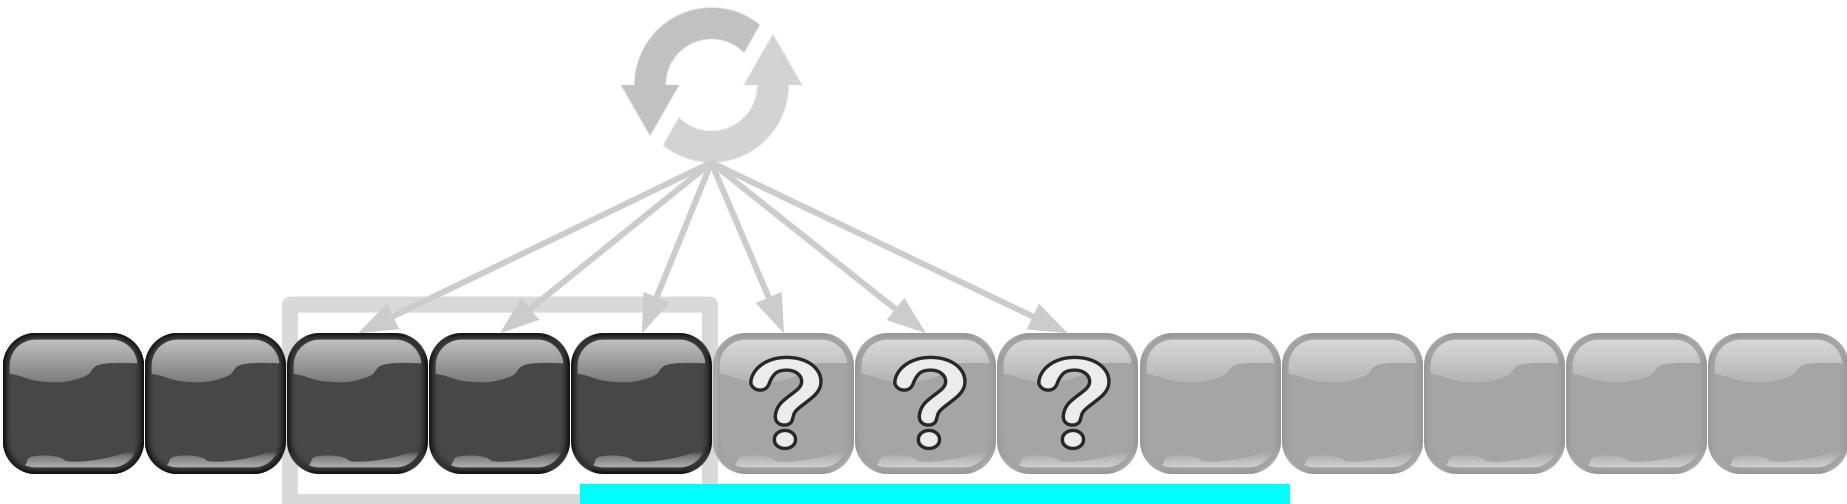






5~8 kitties so far

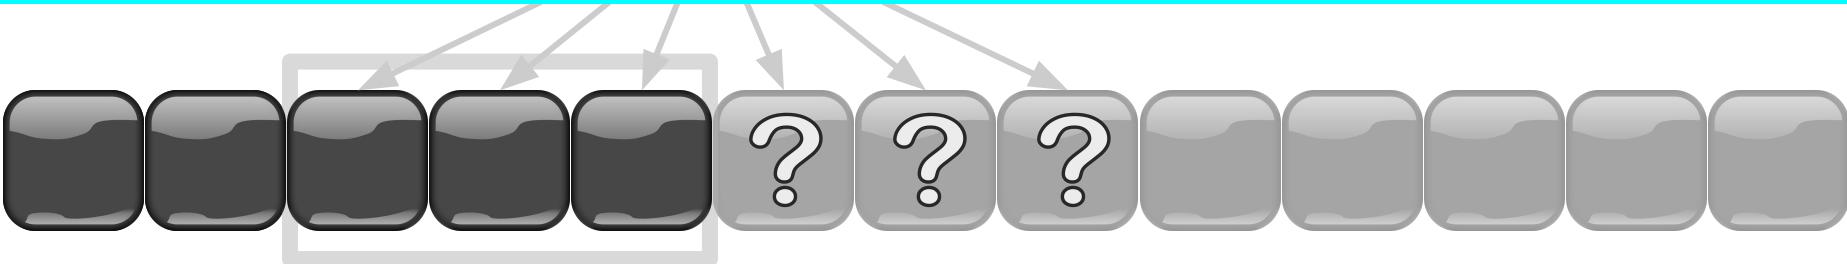




polylog(N) error,
DP estimate



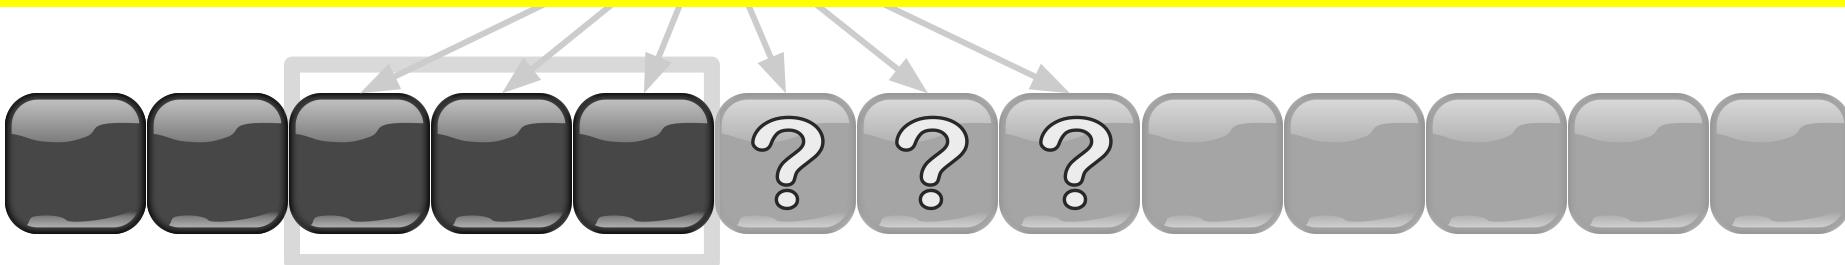
Completes in
O($N \log \log N$)
time





Need:

Oblivious and **DP** alg. that estimates
all prefix sums, with **polylog** error





Naive algorithm:

- Compute all N prefix sums
- Add **independent** noise to each

All prefix sums -- DP and Oblivious

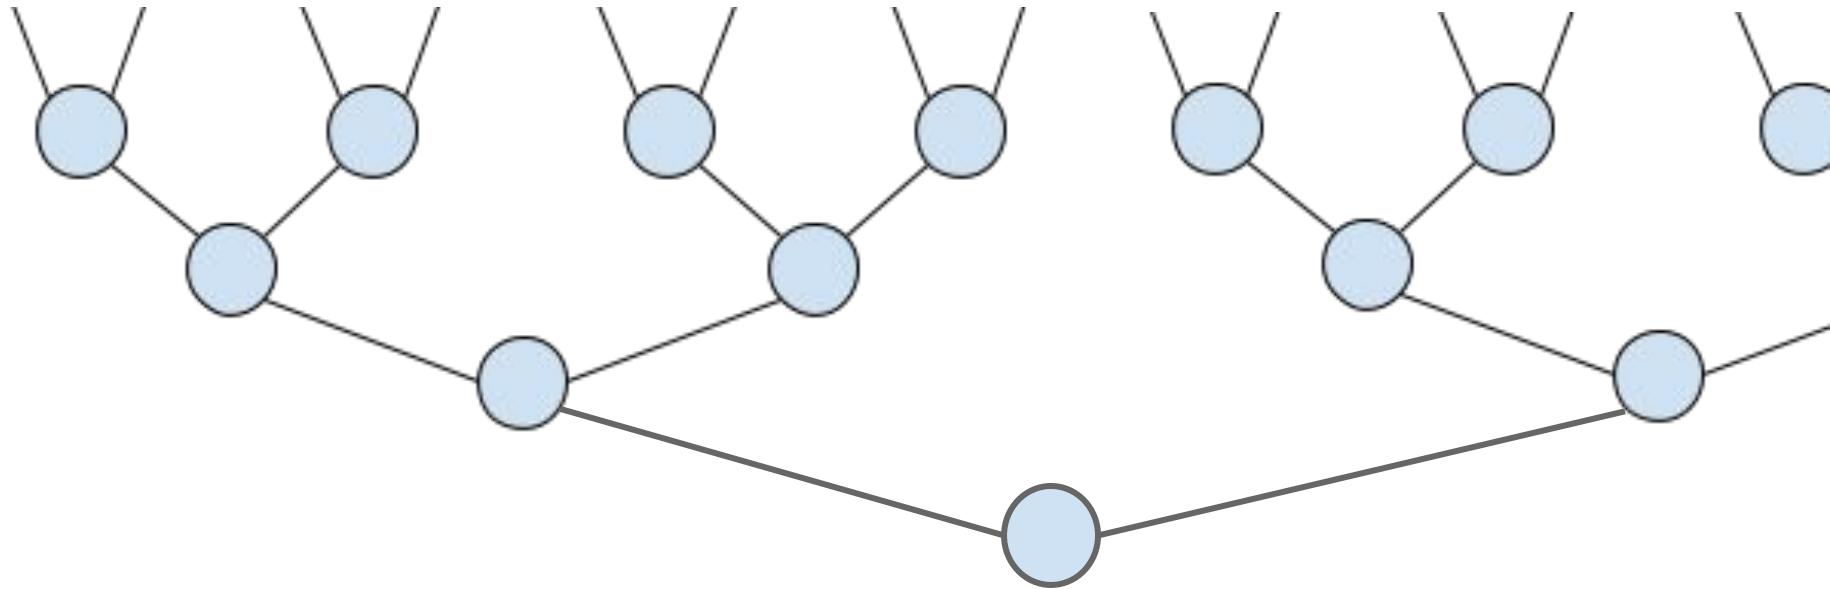


Naive algorithm:

- Compute all N prefix sums
- Add **independent** noise to each

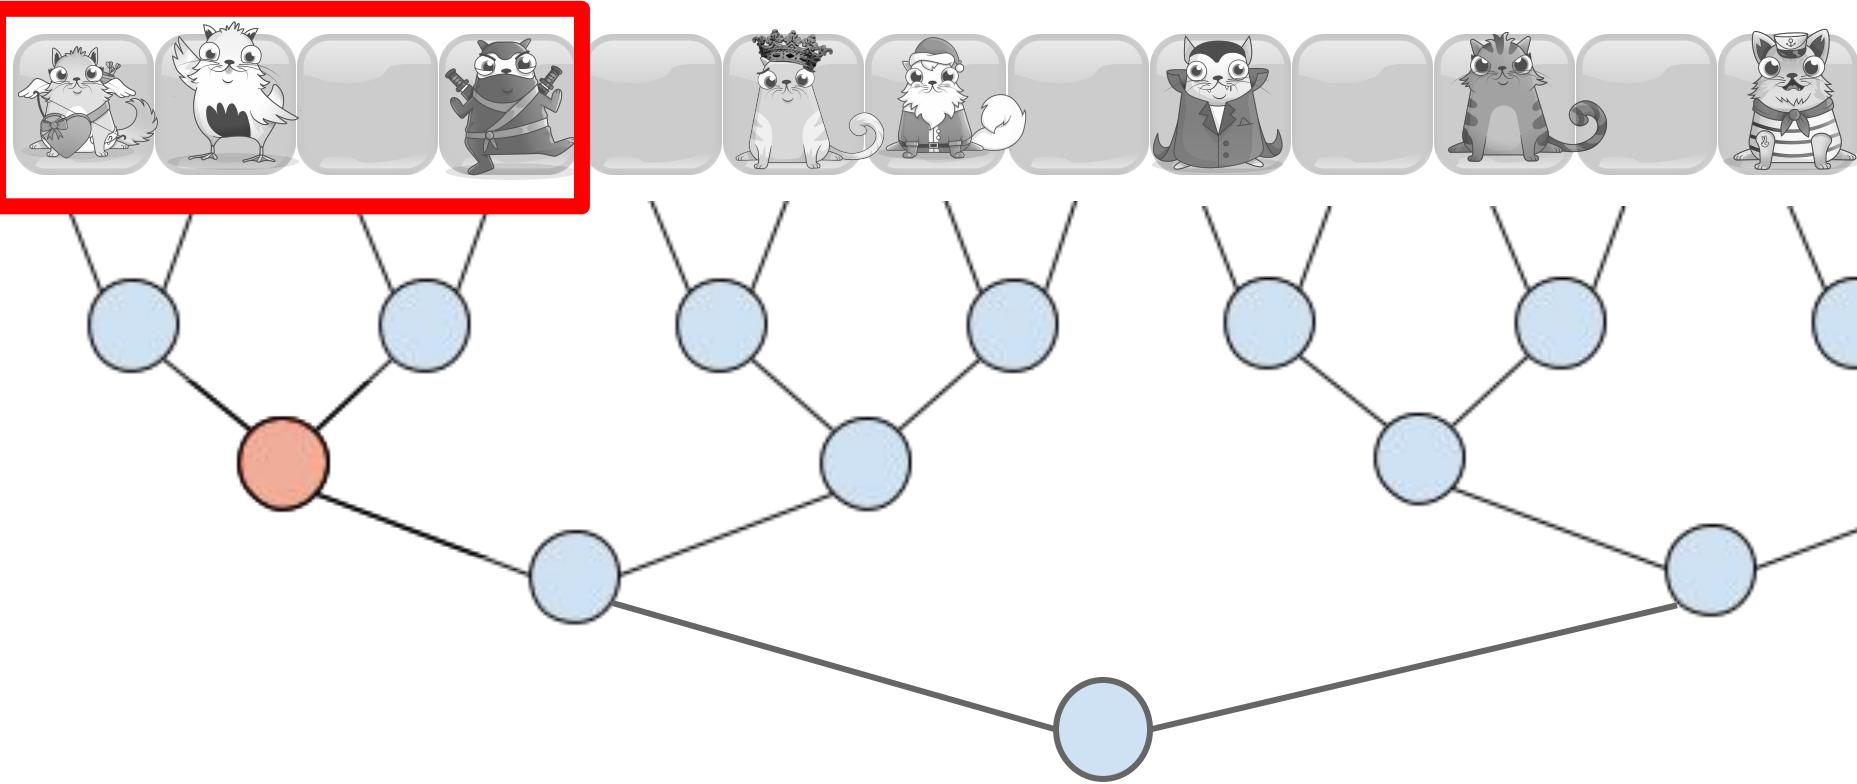


Incurs $\Theta(N)$ error

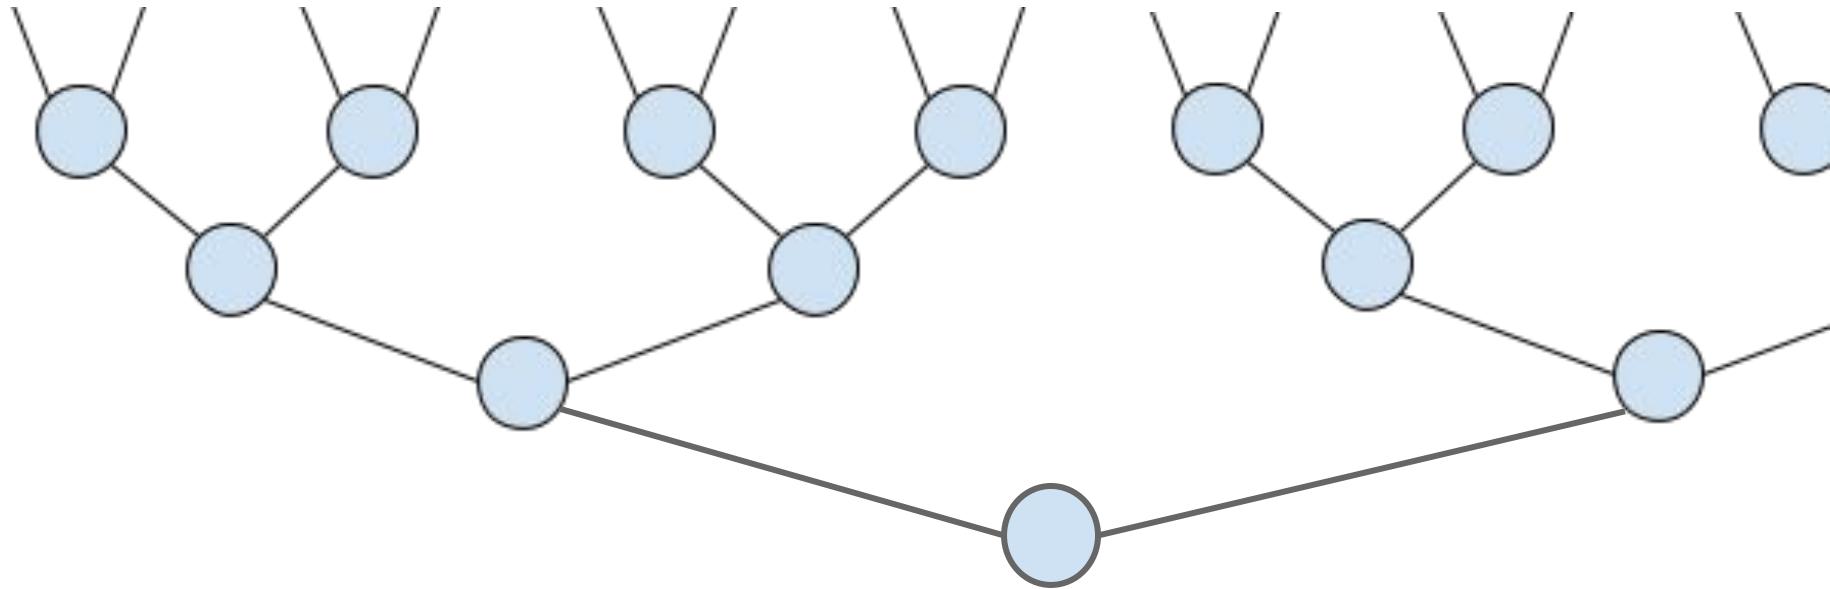


[Dwork et al. 10, C~~S~~'10]

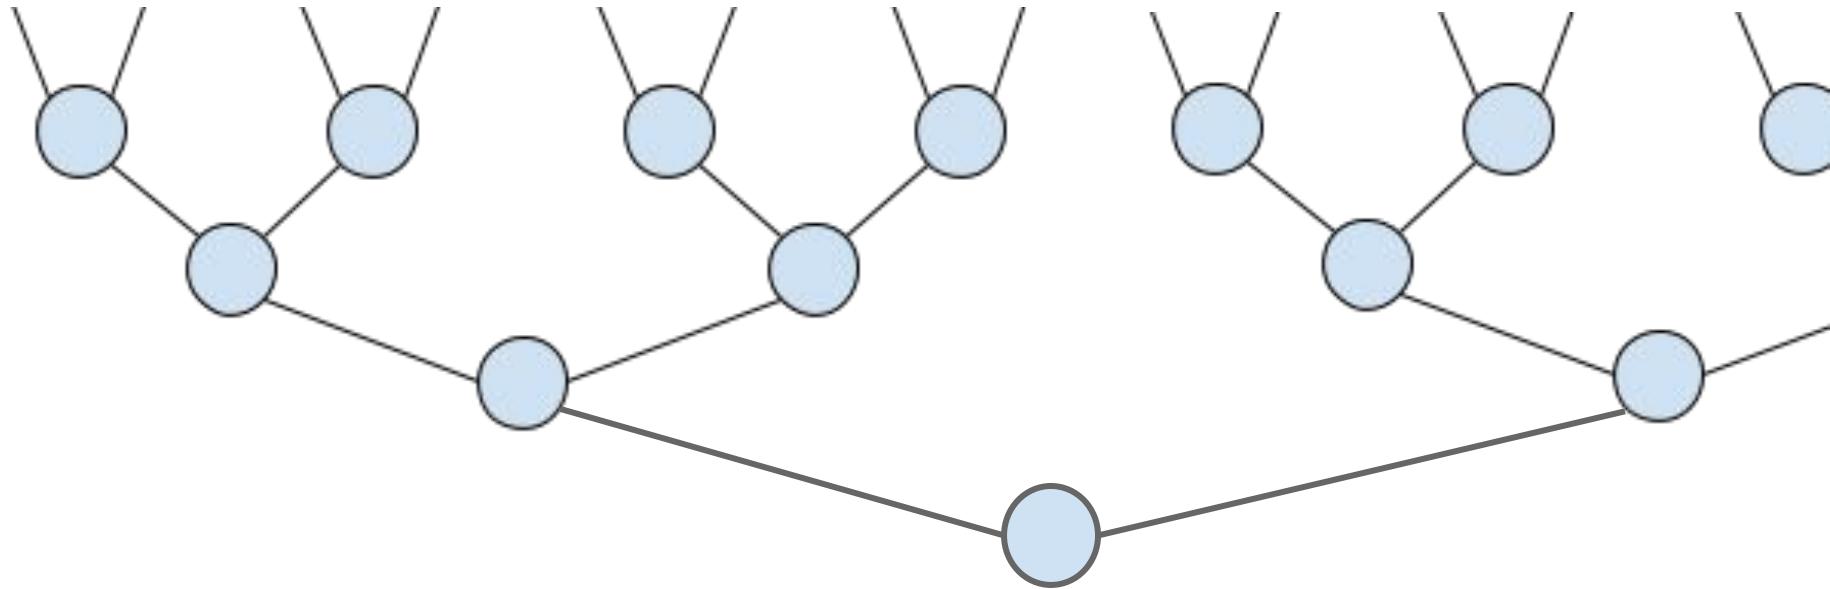
All prefix sums -- DP and Oblivious



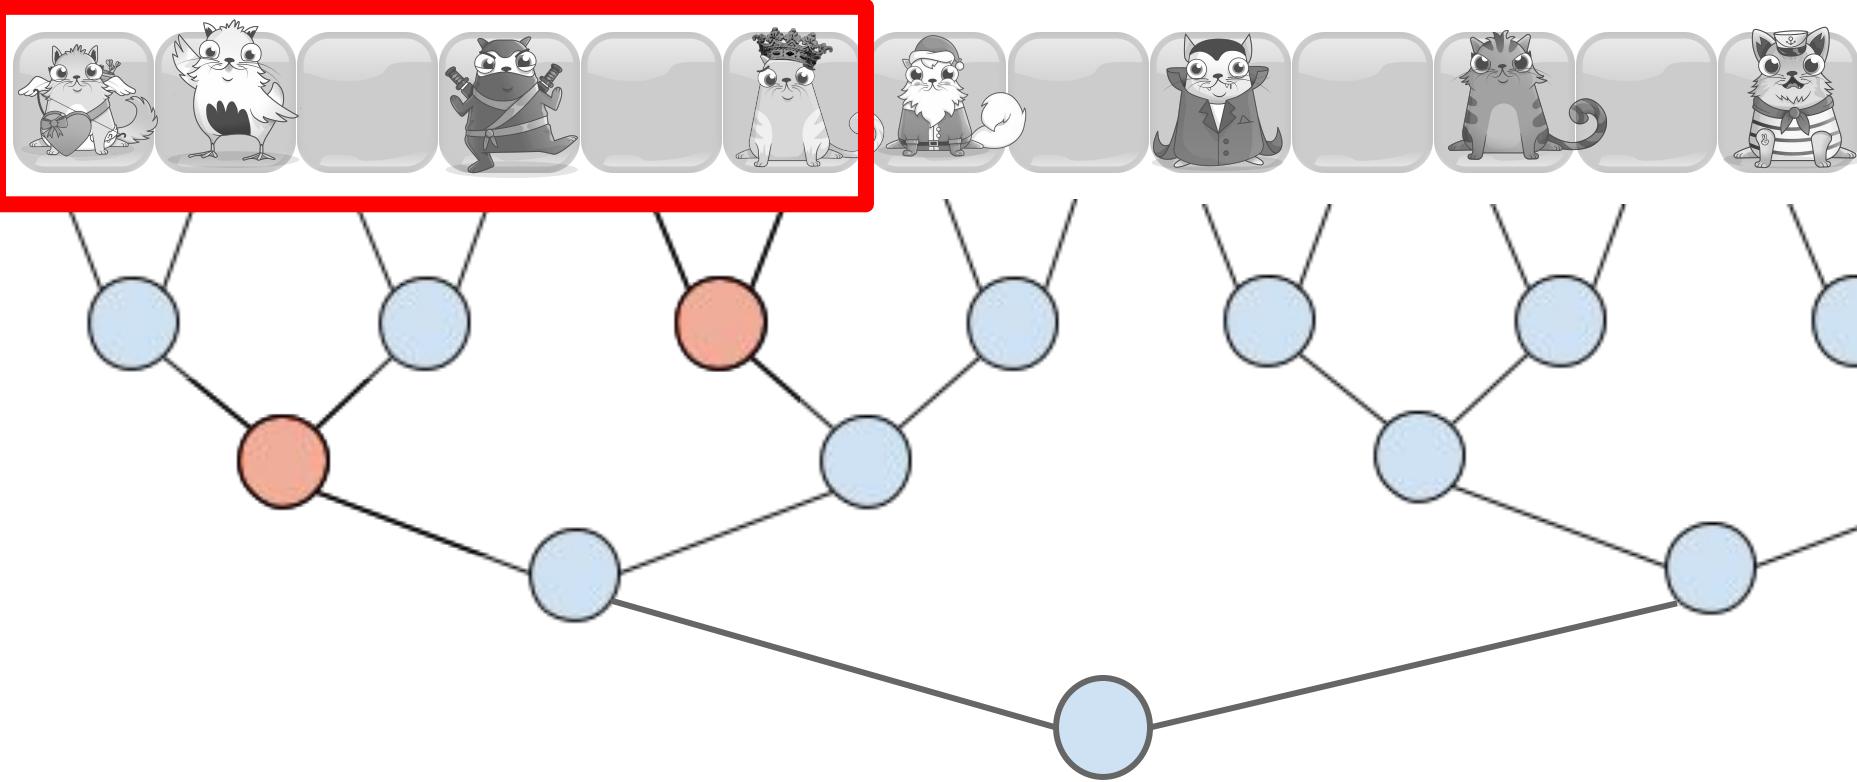
- Every node in the tree represents a range



- Every node in the tree represents a range
- Compute DP estimate for every node in the tree



- Every input appears in only $\log N$ nodes!
- Achieve only $\Theta(\log N)$ error per node!



- Every prefix sum is the sum of $\log N$ nodes
- Achieve $\text{poly log } N$ error for each prefix sum

Summary: Leak rough notion of progress



Non-trivial combination of DP and oblivious algorithms



Apply oblivious alg to small bins



Make DP mechanisms oblivious

Putting it altogether

There exists an $O(N \log \log N)$ time, $(\Theta(1), \text{negl}(N))$ -DO algorithm that realizes stable compaction

Is this necessary?

There exists an $O(N \log \log N)$ time, $(\Theta(1), \text{negl}(N))$ -DO algorithm that realizes stable compaction

$(\epsilon, 0)$ -Differentially Oblivious Stable Compaction:

N log N is necessary

even when ϵ is arbitrarily large!



Other Results in Our Paper

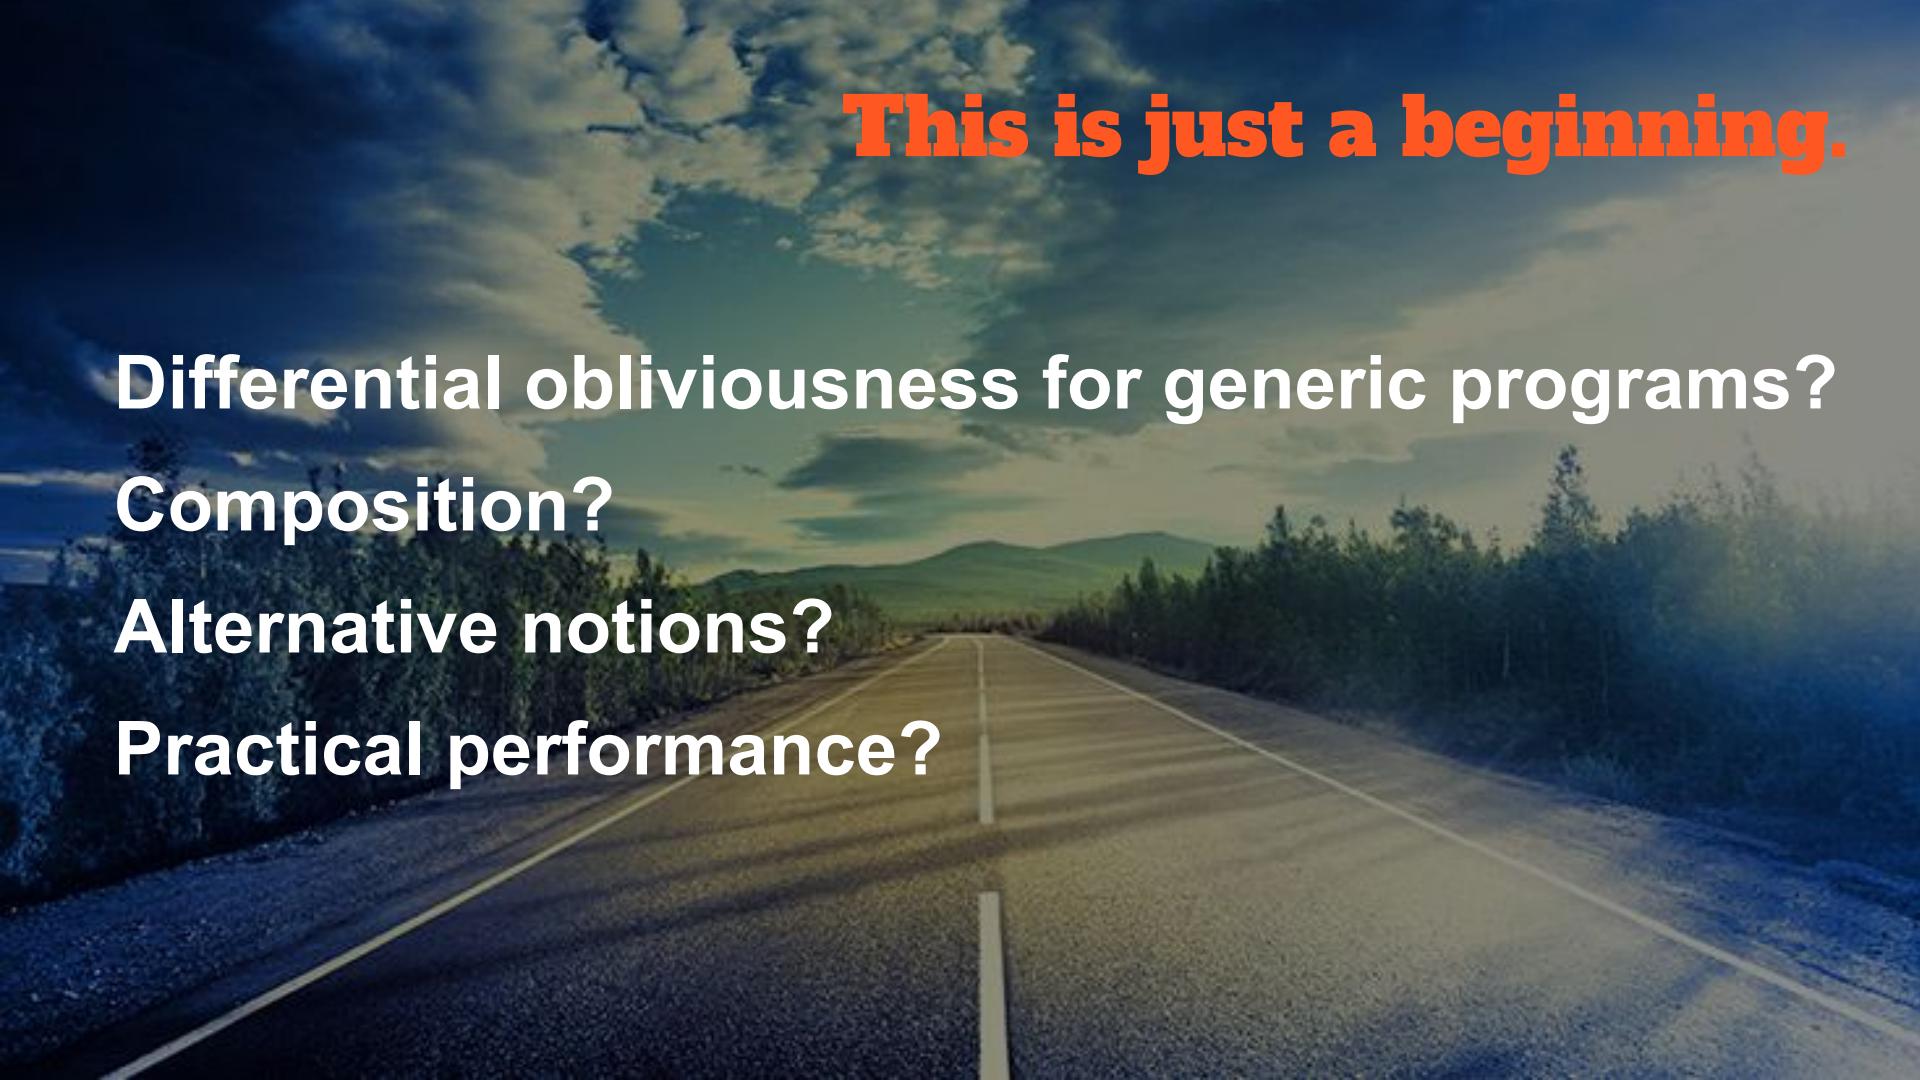
- **lower bounds for obliviousness**
- **merging, range query DB**
 - Differentially oblivious algorithms with $O(\log \log N)$ blowup.
 - $\Omega(\log N)$ blowup necessary for full obliviousness.

Closely Related Works

[Wagh et al.] DP-ORAM, achieve $O(1)$ gain

[Kellaris et al.] DP for length, otherwise fully oblivious

[Mazloom et al.] DP access patterns for MPC

A photograph of a two-lane asphalt road stretching into the distance through a dense green forest. The sky above is filled with large, white, billowing clouds against a blue sky. The perspective of the road creates a sense of depth and movement.

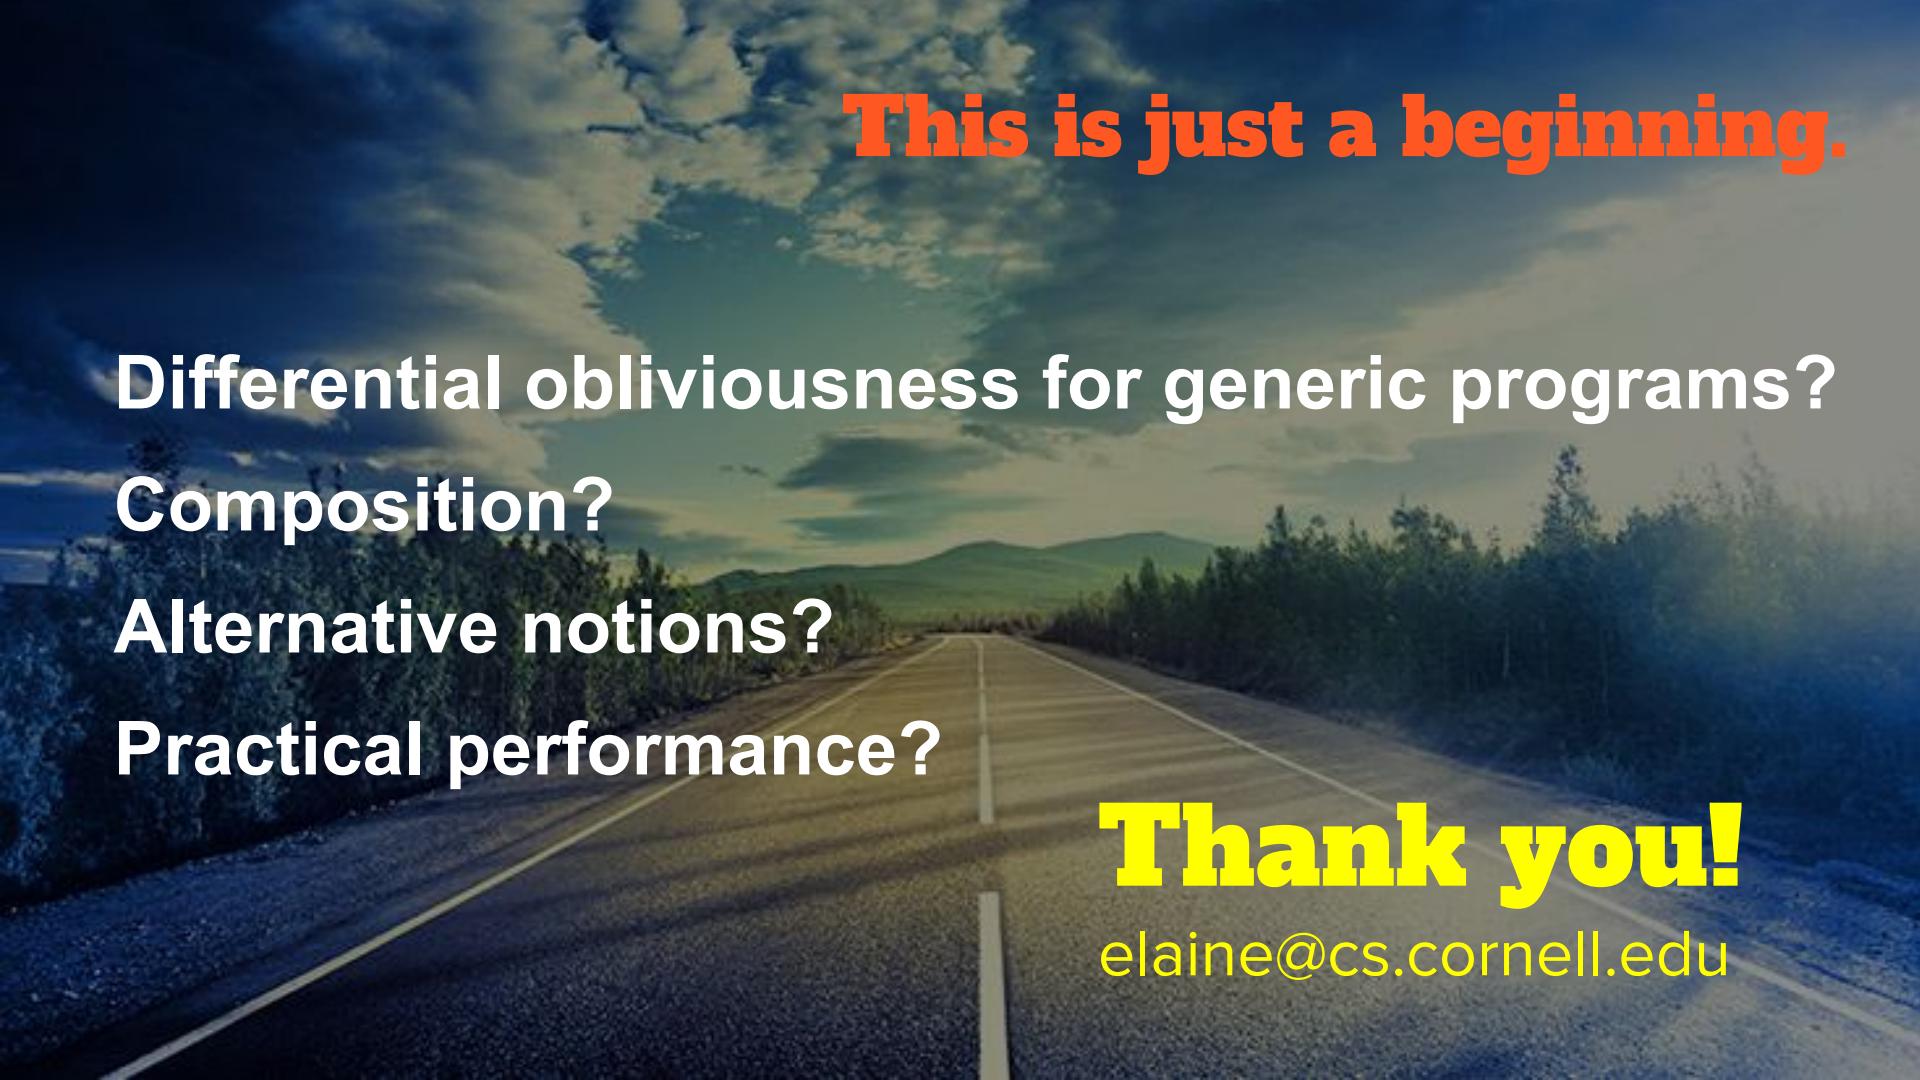
This is just a beginning.

Differential obliviousness for generic programs?

Composition?

Alternative notions?

Practical performance?

A photograph of a two-lane asphalt road stretching into the distance, flanked by dense green trees on both sides. The sky above is filled with large, billowing clouds, creating a dramatic and somewhat mysterious atmosphere.

This is just a beginning.

Differential obliviousness for generic programs?

Composition?

Alternative notions?

Practical performance?

Thank you!

elaine@cs.cornell.edu